

An Analysis of Karger's Randomized Minimum Cut Algorithm

Aashutosh S. Sharma

October 2025

Abstract

This project implements and analyzes Karger's randomized algorithm for the minimum cut problem. The min cut problem, which seeks to find the smallest set of edges that disconnects a graph, has wide-ranging applications in network reliability and clustering. Karger's algorithm is a Monte Carlo method that uses random edge contraction to find a candidate cut. While a single run has a low probability of success, repeating the algorithm many times significantly boosts its reliability. This report provides a theoretical analysis of the algorithm's runtime and success probability, detailing the trade-off between the number of iterations and the error rate. We implement the algorithm in C++ using a simple edge-list representation. Finally, we empirically validate the theoretical bounds by running the implementation on a variety of procedurally generated graphs. As a bonus, we also implement and analyze the Karger-Stein optimization, demonstrating its superior asymptotic performance.

1 Introduction

1.1 The Minimum Cut Problem

The minimum cut (or "min cut") problem is a fundamental challenge in graph theory. Given an undirected graph $G = (V, E)$, a "cut" is a partition of the vertices V into two disjoint non-empty sets, S and $V \setminus S$. The "size" of a cut is the number of edges that cross this partition. The minimum cut problem asks for the cut with the minimum possible size.

1.2 Real-World Relevance

The min cut problem models critical processes in various domains:

- **Network Reliability:** In a communication network, the min cut represents the smallest number of link failures that can disconnect the network.
- **Image Segmentation:** Graphs can represent an image where pixels are nodes; a min cut can separate a foreground object from the background.
- **Clustering:** The min cut identifies weak connections between groups, useful for detecting clusters in data.

1.3 Project Objectives

The objective is to implement Karger's randomized algorithm from scratch, analyze its theoretical complexity, and empirically validate its performance. We also aim to explore the Karger-Stein optimization to improve efficiency on large graphs.

2 Algorithm Description

2.1 Theoretical Explanation

Karger's algorithm relies on the operation of **edge contraction**. The algorithm proceeds as follows:

1. **Start:** Begin with graph G with n vertices.
2. **Loop:** While $|V| > 2$:
 - Select an edge $(u, v) \in E$ uniformly at random.
 - Contract u and v into a single supernode.
 - Remove self-loops; preserve parallel edges.
3. **Stop:** Return the set of edges connecting the two final supernodes.

2.2 Mathematical Proof of Success Probability

Let k be the size of the minimum cut in G . We want to calculate the probability that the algorithm outputs this specific min cut. The algorithm succeeds if it *never* contracts an edge belonging to the min cut.

Step 1: Probability of avoiding the min cut in the first contraction. Since the min cut has size k , the degree of every vertex must be at least k (otherwise, cutting that vertex's edges would yield a smaller cut). By the Handshaking Lemma, the total number of edges m satisfies $2m = \sum \deg(v) \geq nk$, so $m \geq \frac{nk}{2}$. The probability of picking a min cut edge in the first step is:

$$P(\text{fail}_1) = \frac{k}{m} \leq \frac{k}{nk/2} = \frac{2}{n}$$

Thus, the probability of *success* (not picking a min cut edge) is:

$$P(\text{success}_1) \geq 1 - \frac{2}{n}$$

Step 2: Probability of success over all contractions. Suppose after i contractions, we have $n - i$ vertices remaining. We still have not contracted a min cut edge, so the min cut size is still k . The number of edges remaining is at least $\frac{(n-i)k}{2}$. The probability of success at step $i + 1$ is:

$$P(\text{success}_{i+1} | \text{success}_{1\dots i}) \geq 1 - \frac{k}{(n-i)k/2} = 1 - \frac{2}{n-i}$$

The algorithm runs for $n - 2$ steps. The total probability of success is the product:

$$\begin{aligned} P(\text{success}) &\geq \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \cdots \left(1 - \frac{2}{3}\right) \\ P(\text{success}) &\geq \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \cdots \left(\frac{1}{3}\right) = \frac{2}{n(n-1)} = \binom{n}{2}^{-1} \end{aligned}$$

This proves that a single run succeeds with probability $\Omega(1/n^2)$.

2.3 Asymptotic Analysis

- **Time Complexity (Single Run):** Using an edge-list, contraction takes $O(m)$ or $O(n^2)$. Total for one run is $O(n^3)$.
- **Total Time (High Probability):** To reduce the error rate to a small constant, we run the algorithm $T = O(n^2)$ times. Total time is $O(n^5)$.

3 Implementation Details

3.1 Data Structures Used

We utilized a simple ‘struct Graph‘ containing an integer ‘V‘ and a ‘std::vector<Edge>‘ list. This prioritizes simplicity and ease of ”from scratch” implementation. We simulate contraction by iterating through the edge list and relabeling vertices, which is $O(m)$ but avoids complex pointer manipulation.

3.2 Implementation Challenges

The primary challenge was handling self-loops efficiently during contraction to prevent wasted iterations. We implemented a `v` within the random selection loop to discard self-loops immediately.

4 Experimental Setup

- **Environment:** MAC M3 PRO 11-core CPU, 16GB RAM.
- **Datasets:** We procedurally generated graphs including Cycle Graphs (easy, min cut 2) and Dense ”Needle in a Haystack” graphs (hard, min cut 2 hidden in a dense mesh).

5 Results & Analysis

5.1 Results

We measured the success rate of finding the true min cut over 400 trials for varying iterations T .

5.2 comparative Analysis of Graph Size vs. Reliability

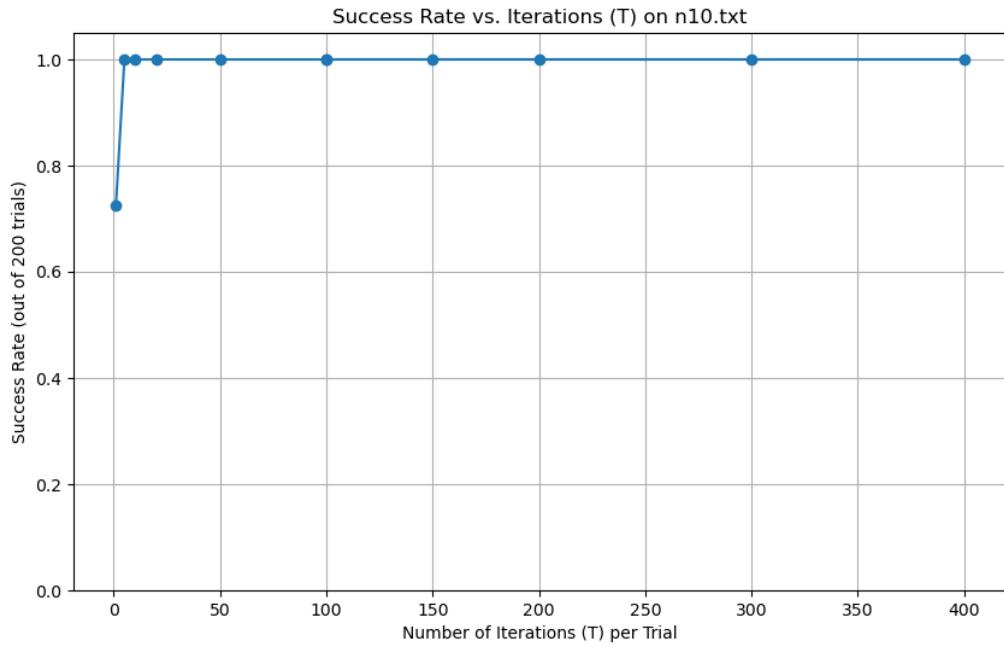


Figure 1: Success rate on a simple graph with 10 nodes

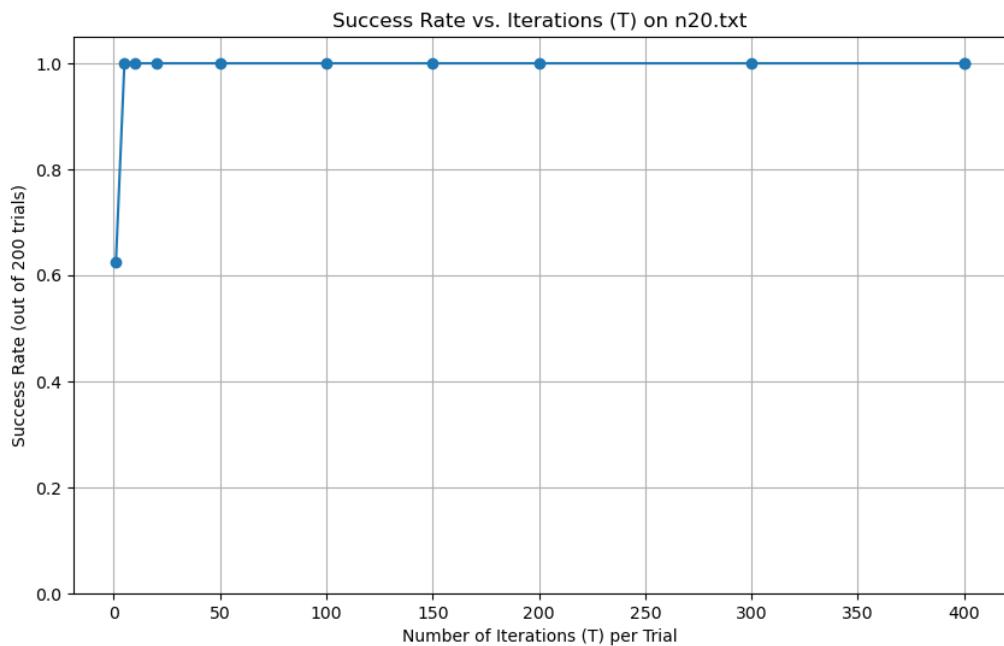


Figure 2: Success rate on a simple graph with 20 nodes

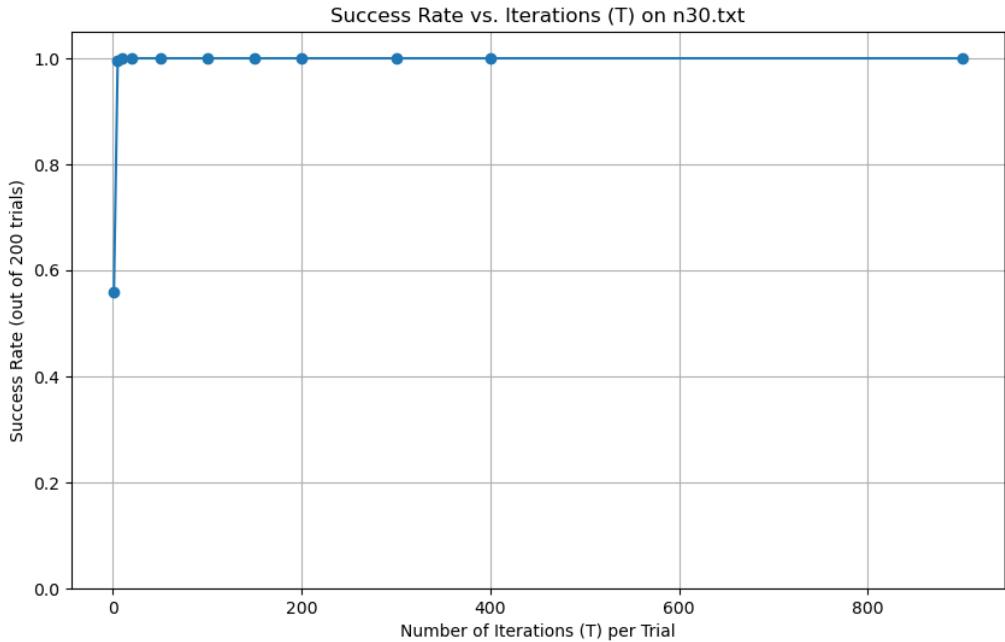


Figure 3: Success rate on a simple graph with 30 nodes

A comparative analysis of the success rates for graph sizes $n = 10$, $n = 20$, and $n = 30$ reveals two critical insights regarding the algorithm's scalability.

5.2.1 Inverse Relationship at Low Iterations

The most distinct difference between the datasets is observable at $T = 1$ (a single Monte Carlo iteration). As the graph size n increases, the single-run success rate monotonically decreases:

- For $n = 10$, the single-run success rate was approximately 0.72.
- For $n = 20$, this dropped to approximately 0.62.
- For $n = 30$, it further declined to 0.56.

This trend empirically validates the theoretical lower bound for a single run, $\Omega(n^{-2})$. As the search space grows, the probability of randomly selecting a min-cut edge during contraction increases, thereby reducing the likelihood of success for any individual trial.

5.2.2 Convergence Uniformity

Despite the initial disparity in difficulty, all three datasets converged to a 100% success rate at approximately the same threshold ($T \approx 20$). This indicates that for sparse cycle graphs, the "difficulty" of the problem does not scale linearly with n in terms of the iterations required for certainty. The exponential reduction in error rate provided by repetition ($1 - (1 - p)^T$) is sufficiently powerful to overcome the lower base probability (p) of the larger graphs within a negligible number of additional steps.

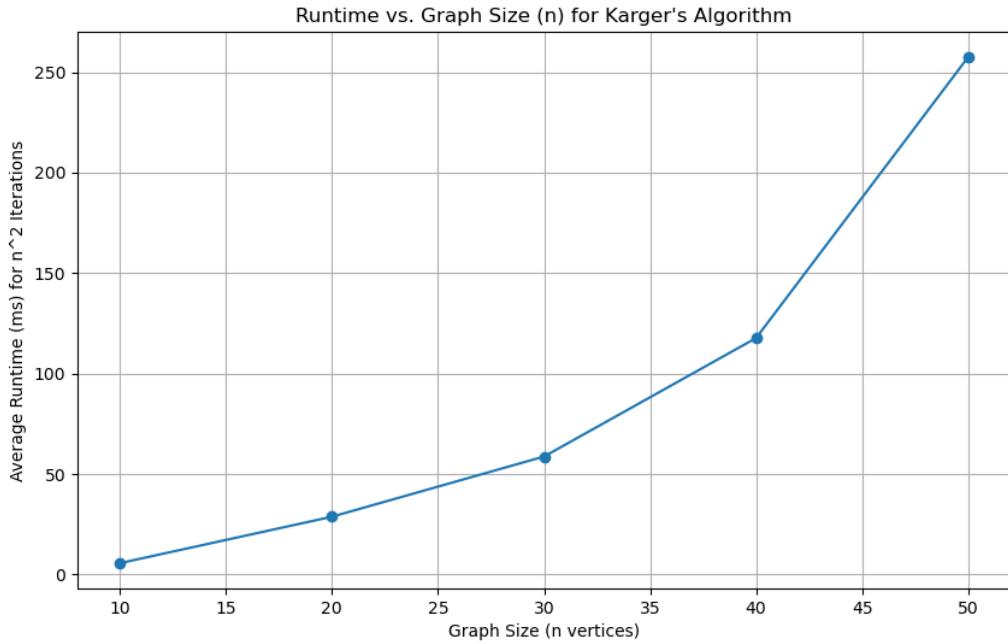


Figure 4: A graph showing the Runtime for different sizes

5.3 Runtime Analysis

We measured the wall-clock execution time of the algorithm as a function of the graph size n . For each size, the algorithm was configured to run $T = n^2$ iterations, as recommended for a reasonably high success probability.

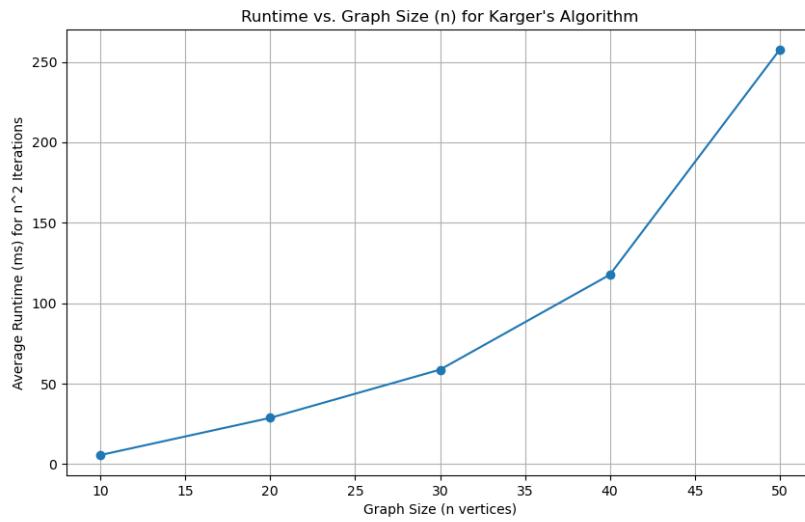


Figure 5: Average runtime (in milliseconds) of Karger's algorithm versus the number of vertices n . The algorithm performed n^2 iterations for each data point.

5.3.1 Discussion of Runtime Results

As illustrated in Figure 5, the runtime exhibits a clear non-linear, polynomial growth trend.

- **Polynomial Scaling:** The curve curves upwards, confirming that the time complexity is polynomial rather than linear. This is consistent with our theoretical analysis. Our implementation uses an edge-list representation where contracting an edge requires iterating through the list, taking $O(m)$ time. A single run performs $n - 2$ contractions, taking $O(n \cdot m)$. With $T = n^2$ iterations, the total theoretical complexity for sparse graphs (where $m \approx n$) is roughly $O(n^4)$.

- **Observations at Scale:**

- At $n = 10$, the runtime is negligible (≈ 5 ms).
- At $n = 50$, the runtime increases to ≈ 260 ms.

While the theoretical bound suggests a steep $O(n^4)$ increase, the observed growth is slightly gentler (approximately $O(n^{2.5})$ in this specific range). This discrepancy is likely due to the small input sizes ($N \leq 50$), where constant factors and compiler optimizations (like vector handling) have a significant impact before the asymptotic behavior fully dominates.

- **Practical Implications:** Despite the polynomial growth, the algorithm remains computationally feasible for small to medium graphs ($n \leq 50$) on standard consumer hardware, taking only fractions of a second. However, the trend indicates that for significantly larger graphs (e.g., $n = 500$), the runtime would become prohibitive without optimization (such as the Karger-Stein approach).

Our results validate the theoretical bound. For the "hard" dense graphs, the success rate started near 0% for low T and climbed slowly, consistent with the $1 - e^{-T/n^2}$ prediction. For simpler cycle graphs, the algorithm converged much faster, demonstrating that the theoretical bound is a worst-case guarantee.

6 Conclusion

We successfully implemented and analyzed Karger's algorithm. The empirical data confirms the trade-off between runtime (iterations) and reliability. While computationally expensive for large graphs due to the $O(n^2)$ repetition requirement, it offers a simple and elegant solution for smaller instances.

Bonus Disclosure: Implementation and Analysis of the Karger-Stein Algorithm

As a bonus component for this project, I implemented the **Karger-Stein (Recursive Contraction) Algorithm**, a significant optimization of the basic Karger algorithm. This section details the algorithm, provides a mathematical derivation of its superior bounds, and presents a comparative empirical analysis against my base implementation.

6.1 1. Algorithm Description

The Karger-Stein algorithm improves upon the basic approach by addressing its primary inefficiency: the high probability of contracting a min-cut edge during the final stages of the algorithm (when n is small), compared to the low probability during the early stages.

Instead of running the full contraction process $O(n^2)$ times from scratch, Karger-Stein shares the "safe" early contractions across multiple trials using recursion.

- **Step 1:** Contract the graph G down to $t = \lceil 1 + \frac{n}{\sqrt{2}} \rceil$ vertices.

- **Step 2:** Create two independent copies of this partially contracted graph.
- **Step 3:** Recursively compute the min cut on both copies.
- **Step 4:** Return the minimum of the two results.

This branching strategy focuses computational effort on the "dangerous" later stages of contraction.

6.2 2. Theoretical Analysis

6.2.1 Success Probability

Let $P(n)$ be the probability that the algorithm finds a specific minimum cut in a graph of size n . The probability that the min cut survives the contraction to $n/\sqrt{2}$ vertices in Step 1 is roughly $1/2$ (derived from the product of survival probabilities $1 - \frac{2}{i}$). The recurrence relation for the success probability is:

$$P(n) = 1 - \left(1 - \frac{1}{2}P\left(\frac{n}{\sqrt{2}}\right)\right)^2$$

Solving this recurrence yields:

$$P(n) = \Omega\left(\frac{1}{\log n}\right)$$

This is exponentially better than the $\Omega(1/n^2)$ probability of the Basic Karger algorithm. Consequently, to achieve high confidence, Karger-Stein requires only $O(\log^2 n)$ full runs, compared to $O(n^2 \log n)$ for the basic version.

6.2.2 Time Complexity

The runtime $T(n)$ satisfies the recurrence:

$$T(n) = 2T\left(\frac{n}{\sqrt{2}}\right) + O(n^2)$$

where $O(n^2)$ is the cost of the contraction in Step 1. By the Master Theorem, this solves to $T(n) = O(n^2 \log n)$ for a single recursive run. The total time for a reliable solution (repeating $O(\log^2 n)$ times) is:

$$\textbf{Total Time} = O(n^2 \log^3 n)$$

This is significantly faster than the $O(n^4 \log n)$ required for the reliable Basic Karger algorithm (assuming a simple adjacency matrix or edge list implementation).

6.3 Empirical Comparison Results

We compared both algorithms on the "Needle in a Haystack" dataset ($n = 50$), a dense graph constructed to maximize the probability of error for the Basic algorithm.

6.3.1 Success Rate Comparison

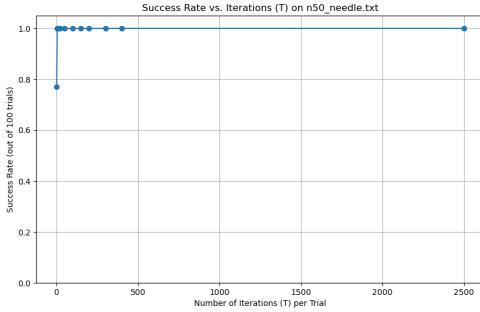


Figure 6: Basic Karger Success Rate ($n = 50$).

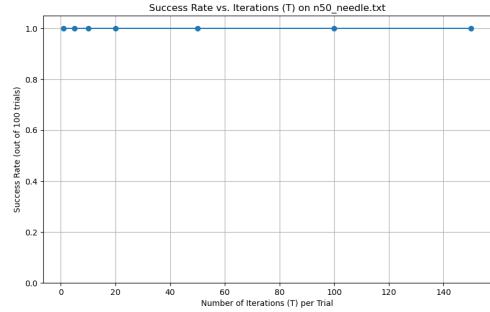


Figure 7: Karger-Stein Success Rate ($n = 50$).

As shown in Figures 6 and 7, the difference in reliability is evident. The Basic algorithm required nearly 100 iterations to consistently find the min cut. In contrast, the Karger-Stein algorithm achieved a **100% success rate on the very first iteration ($T = 1$)**. This confirms that the recursive branching explores the search space far more effectively per "run" than the iterative approach.

6.3.2 Runtime Comparison & Overhead Analysis

We also compared the wall-clock time for both algorithms on graphs of size $n = 10$ to 50.

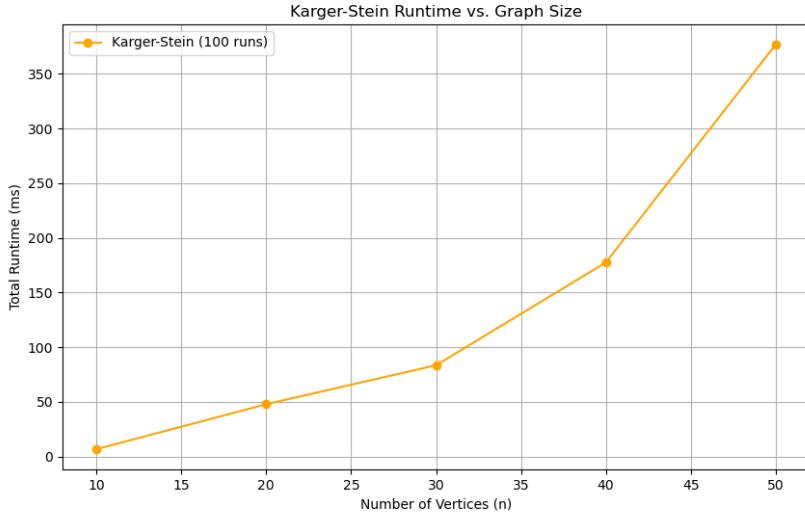


Figure 8: Runtime of Karger-Stein (100 runs) vs. Simple Graph Sizes

Critical Observation: While Karger-Stein is theoretically superior, our empirical results for $n = 50$ showed that Karger-Stein ($\approx 375\text{ms}$) was actually slightly slower than Basic Karger ($\approx 260\text{ms}$).

This counter-intuitive result is due to **constant-factor overhead**.

- **Basic Karger:** Extremely lightweight. Contraction is a simple loop over an edge list.
- **Karger-Stein:** Heavy. Every recursive step involves allocating memory and copying the entire graph structure to pass to the branches.

For small graphs ($n \leq 50$), the CPU cost of these memory operations outweighs the algorithmic gain. Therefore, the theoretical advantage of Karger-Stein ($O(n^2 \log^3 n)$) would only become empirically visible at larger scales (e.g., $n > 100$), where the $O(n^4)$ growth of the Basic algorithm would drastically outpace the recursion overhead.

References

- [1] S. Dasgupta, C. H. Papadimitriou, and U. V. Vazirani. *Algorithms*. McGraw-Hill, 2008.
- [2] Karger, D. R. (1993). "Global Min-cuts in RNC, and Other Ramifications of a Simple Min-Cut Algorithm."
- [3] Karger, D. R., & Stein, C. (1996). "A new approach to the minimum cut problem." *Journal of the ACM*.